

## Exercise: Understanding Language Models

**Goals** The main goal of this worksheet is for you to understand language models a bit more before you implement them.

**Question 1** Consider the prefix “*LeBron James talked about \_*”. Think about the different  $n$ -gram orders here:  $n = 2$  (1 word of context for the language model, just *about*), through  $n = 5$  (all four words of context). **Do you think  $n = 3$  or  $n = 4$  will yield the same distribution over next words as  $n = 5$ ? Why or why not?**

It will be different because you really need to see the whole context to see LeBron. Even “James talked about” leads to a quite distribution of following words (who is James?).

**Question 2** Consider the following corpus (collection of sentences), extended from the one in the video:  
*I like to eat cake but I want to eat pizza right now. Mary told her brother to eat pizza too. He went to Pizza Hut to get some.*

**What is the probability distribution of words following *to* under a 2-gram model?** That is, what is  $P(y \mid \text{to})$ ? Hint: this should be a list of words, each one associated with a probability. You don’t need to explicitly write down all of the words with zero probability.

$P(\text{eat} \mid \text{to}) = \frac{3}{5}$ ,  $P(\text{Pizza} \mid \text{to}) = \frac{1}{5}$ ,  $P(\text{get} \mid \text{to}) = \frac{1}{5}$ , all other words in the vocabulary have 0 probability.

**Question 3** What data structure or data structures would you use to store the words and probabilities for  $P(y \mid \text{to})$ ?

The best answers are two parallel arrays (one for words, one for probabilities) or two parallel ArrayLists/Lists. A Map from String to Double (Java) or dict (Python) is also acceptable.

**(Optional) Question 4** Now suppose you were going to store the entire 2-gram model: the words and probabilities  $P(y \mid x)$  for every  $(x, y)$  pair. What data structure or data structures would you use for this?

This is similar to Question 3 but scaled up. One way to do it is to have two Lists of Lists (or two 2D arrays), where each row stores a distribution  $P(y \mid x)$  for a certain  $x$  following the method in Question 3. However, you still need a way of knowing which row (which index in the outer list) refers to which word  $x$ . The best solution is to have an “indexer,” or a separate list that associates each word with a canonical index. If you have this, you can actually get away with a single 2D array called `probabilities` for everything: you can put  $x$  and  $y$  through the indexer and then access `probabilities[x][y]`.

There are many possible solutions and they get a bit tricky, hence why this part is optional.