

NLP Module: Description for Teachers (UT Austin / Institute for Machine Learning)

Overview This document describes an outreach module intended to introduce high school students to basic concepts from natural language processing (NLP). NLP is a discipline around building systems to solve problems involving human language input, such as question answering, text summarization, machine translation (i.e., Google Translate), and more. Specifically, we discuss language models, including BERT and GPT-3, which power many recent advances at big tech companies and have received lots of attention both in the NLP research community and the media.¹ This module starts from the basics, giving students an understanding of the underlying concepts that power these models via a coding exercise, then uses a web demo² to give them an opportunity to play around with powerful neural network models.

The main model presented that students will actually implement is a *bigram* language model, which models the probability distribution $P(\text{next word} \mid \text{current word})$. For instance, if you see the word *to*, maybe *France* has probability 0.1, *Spain* has probability 0.08, and so on and so forth. Such models can be used for tasks like predictive text on a phone.

We introduce the concepts of basic probability (including conditional probability) and estimating probabilities from data. Learning probabilities for this model does not involve complicated algorithms, but instead just an intuitive counting scheme: if we see *France* as the next word after *to* 10% of the time that *to* shows up in our dataset, then we set $P(\text{France} \mid \text{to}) = 0.1$. Implementing and querying this model largely uses standard math and data structures, but the data structures to store all the quantities needed (particularly counts of neighboring pairs of words) may be a little tricky for students just starting out. There are several levels of programming assignments available depending on students' experience levels.

Goals

1. Expose students to basic concepts from natural language processing and machine learning, including principles of basic probability and learning models on training data
2. Give students a programming challenge involving querying data structures, basic loop constructs, and mathematical operations
3. Give students hands-on experience with state-of-the-art language models via a web demo
4. Challenge students to look at language data and think critically about what the pros and cons of these models are
5. Give students pointers to other resources to follow up on these topics

What this module contains

10 videos (approx. 80 minutes), 3 exercises (60+ minutes)

- Introductory content about NLP, machine learning, and high-level ideas of language modeling (3 videos, 19 minutes)
- In-depth about how to implement n -gram language models (3 video, 28 mins, including about 10 minutes of code walkthrough) (**exercise 0**; 15 minutes)

¹These have been covered extensively in the media; e.g., this article about GPT-2, the predecessor to GPT-3, when it came out: <https://www.cnn.com/2019/02/18/tech/dangerous-ai-text-generator>

²<https://transformer.huggingface.co/doc/distil-gpt2>

- Hands-on exercise with n -gram language models (**exercise 1**; 20 minutes - 2 hours depending on what version)
- Introduction to Write With Transformer (8 minute video)
- Hands-on exercise with web demo and discussion (5 minute video + **exercise 2**, 10-15 minute)
- Broader applications and where to go next (2 videos, 12 minutes)

Video titles

1. What is NLP?
2. Machine Learning
3. Language Modeling (exercise 0)
4. Building n -gram LMs (exercise 0)
5. Bigram LM Code (exercise 1) **note: there are versions of this video for both Java and Python**
6. Querying the LM (exercise 1) **note: there are versions of this video for both Java and Python**
7. Hands-on: Write With Transformer (exercise 2)
8. Discussion: Write With Transformer (exercise 2)
9. Language Models in the News
10. What's Next?

Exercises

- Exercise 0: Pen-and-paper exercises around understanding language models and the data structures involved with them. Solutions can be found at <https://cs.utexas.edu/~gdurrett/courses/nlp-module/exercise0-solutions.pdf>
- Exercise 1: Implementing bigram language models. Students will use basic data structures and provided framework code to query a probabilistic model of what the next word in a sentence is likely to be. Framework code reads in data and populates the needed data structures; students primarily implement loop logic to sample sentences from the probability distribution placed by the model. This exercise follows a worksheet.
- Exercise 2: Students play around with the Write With Transformer demo (see URL in footnote on page 1). The video lectures include several questions which students will answer using the web tool.

How to use this module

This module can be tailored to contain different amounts of coding and be suitable for students at different experience levels.

We recommend using this module in AP Computer Science A. It fits most naturally into Unit 7, Topics 7.1-7.3.³ The two coding parts feature incremental construction of an ArrayList and iteration through an ArrayList with a stopping condition based on the contents. Exercise 0 also has a conceptual question about 2D arrays (Unit 8), but you can skip this.

Quick / non-coding version You can use **videos 1-4 and 7-10** as an introduction to the concepts that doesn't involve coding. This plus the hands-on Exercise 2 can be run for an audience with little CS background, although there are some mathematical concepts that students may not have seen.

Choices for the classroom

Alternative order: Put exercise 2 first Exercise 1 involves coding and Exercise 2 is just a web demo. You may wish to use the videos in the order: 1-4 + Exercise 0, 7-8 + Exercise 2, 5-6 + Exercise 1. If you think your class would benefit from the hands-on exercise first before diving in, you can structure it in this way.

Points to pause You may wish to pause more frequently in the videos than I do. Some additional recommended pause points:

Video 1: after any of the example tasks if you wish to discuss them

Video 2: timestamp 3:32 (right after Cleverbot, if you wish to play with it)

Video 3: timestamp 3:00 (check comprehension after “these have to add up to 1”), timestamp 10:00 (after the slide “N-gram LM” is complete)

Video 6: after the code demonstration is complete when we switch back to the slide presentation

Code and Platforms

This module supports both Java and Python versions; the two videos that depend on code specifics have been recorded twice, once for each language. Code for both versions is available either as stand-alone downloads or as web-based projects through `repl.it`.

³<https://apcentral.collegeboard.org/media/pdf/ap-computer-science-a-course-and-exam-description.pdf> at the time of this writing in August 2022