

Greg Durrett
gdurrett@nyu.edu
<https://gregdurrett.github.io>

POSITIONS

New York University

Associate Professor, Department of Computer Science (Courant Institute) & Center for Data Science, September 2025 –

The University of Texas at Austin

Associate Professor, Department of Computer Science, September 2023 – August 2025

Assistant Professor, Department of Computer Science, September 2017 – August 2023

EDUCATIONAL HISTORY

University of California, Berkeley

Ph.D. in Computer Science, August 2016

Thesis: *Identifying and Resolving Entities in Text* (Advisor: Dan Klein)

Massachusetts Institute of Technology

B.S. in Computer Science and B.S. in Mathematics, minor in Music, June 2010

PUBLICATIONS

Publications and recent preprints can be found on <https://taur.cs.utexas.edu/>

109. Manya Wadhwa, Zayne Sprague, Chaitanya Malaviya, Philippe Laban, Junyi Jessie Li, and **Greg Durrett**. 2025. EvalAgent: Discovering Implicit Evaluation Criteria from the Web Paper. *Proceedings of the Conference on Language Modeling (COLM)*.
108. Juan Diego Rodriguez, Wenxuan Ding, Katrin Erk, and **Greg Durrett**. 2025. RankAlign: A Ranking View of the Generator-Validator Gap in Large Language Models. *Proceedings of the Conference on Language Modeling (COLM)*.
107. Anirudh Khatry, Robert Zhang, Jia Pan, Ziteng Wang, Qiaochu Chen, **Greg Durrett**, and Isil Dillig. 2025. CRUST-Bench: A Comprehensive Benchmark for C-to-safe-Rust Transpilation. *Proceedings of the Conference on Language Modeling (COLM)*.
106. Ramya Namuduri, Yating Wu, Anshun Asher Zheng, Manya Wadhwa, **Greg Durrett**, and Junyi Jessie Li. 2025. QUDSim: Quantifying Discourse Similarities in LLM-Generated Text. *Proceedings of the Conference on Language Modeling (COLM)*.
105. Xi Ye, Fangcong Yin, Yinghui He, Joie Zhang, Howard Yen, Tianyu Gao, **Greg Durrett**, and Danqi Chen. Long-Proc: Benchmarking Long-Context Language Models on Long Procedural Generation. *Proceedings of the Conference on Language Modeling (COLM)*.
104. Tuhina Tripathi, Manya Wadhwa, **Greg Durrett**, and Scott Niekum. Pairwise or Pointwise? Evaluating Feedback Protocols for Bias in LLM-Based Evaluation. *Proceedings of the Conference on Language Modeling (COLM)*.
103. Mahnaz Koupaei, Xueying Bai, Mudan Chen, **Greg Durrett**, Nathanael Chambers, and Niranjan Balasubramanian. 2025. Causal Graph based Event Reasoning using Semantic Relation Experts. *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*.
102. Xinyu Zhao, Fangcong Yin, and **Greg Durrett**. 2025. Understanding Synthetic Context Extension via Retrieval Heads. *Proceedings of the International Conference on Machine Learning (ICML)*.
101. Zayne Sprague, Fangcong Yin, Juan Diego Rodriguez, Dongwei Jiang, Manya Wadhwa, Prasann Singhal, Xinyu Zhao, Xi Ye, Kyle Mahowald, and **Greg Durrett**. 2025. To CoT or not to CoT? Chain-of-thought helps mainly on math and symbolic reasoning. *Proceedings of the International Conference on Learning Representations (ICLR)*.
100. Thom Lake, Eunsol Choi, and **Greg Durrett**. 2025. From Distributional to Overton Pluralism: Investigating Large Language Model Alignment. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
99. Fangcong Yin, Xi Ye, and **Greg Durrett**. 2024. LoFiT: Localized Fine-tuning on LLM Representations. *Advances in Neural Information Processing Systems (NeurIPS)*, **spotlight presentation at MINT workshop**.

98. Liyan Tang, Philippe Laban, and **Greg Durrett**. 2024. MiniCheck: Efficient Fact-Checking of LLMs on Grounding Documents. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
97. Abhishek Divekar and **Greg Durrett**. 2024. SynthesizRR: Generating Diverse Datasets with Retrieval Augmentation. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
96. Yating Wu, Ritika Mangla, Alexandros Dimakis, **Greg Durrett**, and Junyi Jessy Li. 2024. Which questions should I answer? Salience Prediction of Inquisitive Questions. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*, *outstanding paper award*.
95. Manya Wadhwa, Xinyu Zhao, Junyi Jessy Li, and **Greg Durrett**. 2024. Learning to Refine with Fine-Grained Natural Language Feedback. *Findings of the Conference on Empirical Methods for Natural Language Processing (EMNLP-Findings)*.
94. Anisha Gunjal and **Greg Durrett**. 2024. Molecular Facts: Desiderata for Decontextualization in LLM Fact Verification. *Findings of the Conference on Empirical Methods for Natural Language Processing (EMNLP-Findings)*.
93. Prasann Singhal, Nathan Lambert, Scott Niekum, Tanya Goyal, and **Greg Durrett**. 2024. D2PO: Discriminator-Guided DPO with Response Evaluation Models. *Proceedings of the Conference on Language Modeling (COLM)*.
92. Prasann Singhal, Tanya Goyal, Jiacheng Xu, and **Greg Durrett**. 2024. A Long Way To Go: Investigating Length Correlations in RLHF. *Proceedings of the Conference on Language Modeling (COLM)*, *oral spotlight (22/299 accepted papers)*.
91. Manya Wadhwa, Jifan Chen, Junyi Jessy Li, and **Greg Durrett**. 2024. Using Natural Language Explanations to Rescale Human Judgments. *Proceedings of the Conference on Language Modeling (COLM)*.
90. Juan Diego Rodriguez, Katrin Erk, and **Greg Durrett**. 2024. X-PARADE: Cross-Lingual Textual Entailment and Information Divergence across Paragraphs. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
89. Jifan Chen, Grace Kim, Aniruddh Sriram, **Greg Durrett**, and Eunsol Choi. 2024. Complex Claim Verification with Evidence Retrieved in the Wild. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
88. Zayne Sprague, Xi Ye, Kaj Bostrom, Swarat Chaudhuri, and **Greg Durrett**. 2024. MuSR: Testing the Limits of Chain-of-thought with Multistep Soft Reasoning. *Proceedings of the International Conference on Learning Representations (ICLR)*, *spotlight*.
87. Jiayi Wei, **Greg Durrett**, and Isil Dillig. 2024. Coeditor: Leveraging Contextual Changes for Multi-round Code Auto-editing. *Proceedings of the International Conference on Learning Representations (ICLR)*, *spotlight*.
86. Xi Ye, Qiaochu Chen, Isil Dillig, and **Greg Durrett**. 2023. SatLM: Satisfiability-Aided Language Models Using Declarative Prompting. *Advances in Neural Information Processing Systems (NeurIPS)*.
85. Shankar Padmanabhan, Yasumasa Onoe, Michael J.Q. Zhang, **Greg Durrett**, and Eunsol Choi. 2023. Propagating Knowledge Updates to LMs Through Distillation. *Advances in Neural Information Processing Systems (NeurIPS)*.
84. Ryo Kamoi, Tanya Goyal, Juan Diego Rodriguez, and **Greg Durrett**. 2023. WICE: Real-World Entailment for Claims in Wikipedia. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
83. Xi Ye and **Greg Durrett**. 2023. Explanation Selection Using Unlabeled Data for Chain-of-Thought Prompting. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
82. Yating Wu, Ritika Mangla, **Greg Durrett**, and Junyi Jessy Li. 2023. QUDeval: The Evaluation of Questions Under Discussion Discourse Parsing. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
81. Jarad Forristal, Fatemehsadat Mireshghallah, **Greg Durrett**, and Taylor Berg-Kirkpatrick. 2023. A Block Metropolis-Hastings Sampler for Controllable Energy-based Text Generation. *Proceedings of the Conference on Natural Language Learning (CoNLL)*.
80. Jocelyn Chen, Arko Banerjee, Çağatay Demiralp, **Greg Durrett**, and Isil Dillig. 2023. Data Extraction via Semantic Regular Expression Synthesis. *Object-oriented Programming, Systems, Languages, and Applications (OOPSLA)*.
79. Liyan Tang, Zhaoyi Sun, Betina Idnay, Jordan G Nestor, Ali Soroush, Pierre A. Elias, Ziyang Xu, Ying Ding, **Greg Durrett**, Justin Rousseau, Chunhua Weng, and Yifan Peng. 2023. Evaluating Large Language Models on Medical Evidence Summarization. *npj Digital Medicine*

78. Yasumasa Onoe, Michael J.Q. Zhang, Shankar Padmanabhan, **Greg Durrett**, and Eunsol Choi. 2023. Can LMs Learn New Entities from Descriptions? Challenges in Propagating Injected Knowledge. *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*.
77. Prasann Singhal, Jiacheng Xu, Xi Ye, and **Greg Durrett**. 2023. EEL: Efficiently Encoding Lattices for Reranking. *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*.
76. Liyan Tang, Tanya Goyal, Alexander R. Fabbri, Philippe Laban, Jiacheng Xu, Semih Yavuz, Wojciech Kryściński, Justin F. Rousseau, and **Greg Durrett**. 2023. Understanding Factual Errors in Summarization: Errors, Summarizers, Datasets, Error Detectors. *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*.
75. Xi Ye, Srinivasan Iyer, Asli Celikyilmaz, Ves Stoyanov, **Greg Durrett**, and Ramakanth Pasunuru. 2023. Complementary Explanations for Effective In-Context Learning. *Findings of Association for Computational Linguistics (Findings of ACL)*.
74. Adithya Bhaskar, Alexander R. Fabbri, and Greg Durrett. Prompted Opinion Summarization with GPT-3.5. *Findings of Association for Computational Linguistics (Findings of ACL)*.
73. Wei-Jen Ko, Yating Wu, Cutter Dalton, Dananjay Srinivas, **Greg Durrett**, and Junyi Jessy Li. 2023. Discourse Analysis via Questions and Answers: Parsing Dependency Structures of Questions Under Discussion. *Findings of Association for Computational Linguistics (Findings of ACL)*.
72. Ryo Kamoi, Tanya Goyal, and **Greg Durrett**. 2023. Shortcomings of Question Answering Based Factuality Frameworks for Error Localization. *Proceedings of the European Chapter of the Association for Computational Linguistics (EACL)*.
71. Prasann Singhal, Jarad Forristal, Xi Ye, and **Greg Durrett**. 2023. Assessing Out-of-Domain Language Model Performance from Few Examples. *Proceedings of the European Chapter of the Association for Computational Linguistics (EACL)*.
70. Mahnaz Koupaei, **Greg Durrett**, Nate Chambers, and Niranjan Balasubramanian. 2023. Modeling Complex Event Scenarios via Simple Entity-focused Questions. *Proceedings of the European Chapter of the Association for Computational Linguistics (EACL)*.
69. Jiayi Wei, **Greg Durrett**, and Isil Dillig. 2023. TypeT5: Seq2seq Type Inference using Static Analysis. *Proceedings of the International Conference on Learning Representations (ICLR)*.
68. Jifan Chen, Aniruddh Sriram, Eunsol Choi, and **Greg Durrett**. 2022. Generating Literal and Implied Subquestions to Fact-check Complex Claims. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
67. Tanya Goyal, Junyi Jessy Li, and **Greg Durrett**. 2022. SNaC: Coherence Error Detection for Narrative Summarization. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
66. Zayne Sprague, Kaj Bostrom, Swarat Chaudhuri, and **Greg Durrett**. 2022. Natural Language Deduction with Incomplete Information. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
65. Kaj Bostrom, Zayne Sprague, Swarat Chaudhuri, and **Greg Durrett**. 2022. Natural Language Deduction through Search over Statement Compositions. *Findings of the Conference on Empirical Methods for Natural Language Processing (EMNLP-Findings)*.
64. Wei-Jen Ko, Cutter Dalton, Mark Simmons, Eliza Fisher, **Greg Durrett**, and Junyi Jessy Li. 2022. Discourse Comprehension: A Question Answering Framework to Represent Sentence Connections. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
63. Xi Ye and **Greg Durrett**. 2022. The Unreliability of Explanations in Few-shot Prompting for Textual Reasoning. *Advances in Neural Information Processing Systems (NeurIPS)*.
62. Qiaochu Chen, Shankara Pailoor, Celeste Barnaby, Abby Criswell, Chenglong Wang, **Greg Durrett**, and Isil Dillig. 2022. Type-Directed Synthesis of Visualizations from Natural Language Queries. *Object-oriented Programming, Systems, Languages, and Applications (OOPSLA)*.
61. Jiacheng Xu, Siddhartha Reddy Jonnalagadda, and **Greg Durrett**. 2022. Massive-scale Decoding for Text Generation using Lattices. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
60. Yasumasa Onoe, Michael J.Q. Zhang, Eunsol Choi, and **Greg Durrett**. 2022. Entity Cloze By Date: Understanding what LMs know about unseen entities. *Findings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL-Findings)*.

59. Benjamin Mariano, Yanju Chen, Yu Feng, **Greg Durrett**, and Isil Dillig. 2022. Automated Transpilation of Imperative to Functional Code using Neural-Guided Program Synthesis. *Object-oriented Programming, Systems, Languages, and Applications (OOPSLA)*
58. Xi Ye and **Greg Durrett**. 2022. Can Explanations Be Useful for Calibrating Black Box Models? *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*.
57. Ojas Ahuja, Jiacheng Xu, Akshay Gupta, Kevin Horecka, and **Greg Durrett**. 2022. ASPECTNEWS: Aspect-Oriented Summarization of News Documents. *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*.
56. Tanya Goyal, Jiacheng Xu, Junyi Jessy Li, and **Greg Durrett**. 2022. Training Dynamics for Text Summarization Models. *Findings of Association for Computational Linguistics (Findings of ACL)*.
55. Yasumasa Onoe, Michael J.Q. Zhang, Eunsol Choi, and **Greg Durrett**. 2021. CREAK: A Dataset for Commonsense Reasoning over Entity Knowledge. *Proceedings of NeurIPS (Datasets and Benchmarks Track)*.
54. Kaj Bostrom, Xinyu Zhao, Swarat Chaudhuri, and **Greg Durrett**. 2021. Flexible Generation of Natural Language Deductions. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
53. Xi Ye, Rohan Nair, and **Greg Durrett**. 2021. Connecting Attributions and QA Model Behavior on Realistic Counterfactuals. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
52. Jifan Chen, Eunsol Choi, and **Greg Durrett**. 2021. Can NLI Models Verify QA Systems’ Predictions? *Findings of the Conference on Empirical Methods for Natural Language Processing (Findings of EMNLP)*.
51. Xi Ye, Qiaochu Chen, Isil Dillig, and **Greg Durrett**. 2021. Optimal Neural Program Synthesis from Multimodal Specifications. *Findings of the Conference on Empirical Methods for Natural Language Processing (Findings of EMNLP)*.
50. Yasumasa Onoe, Michael Boratko, Andrew McCallum, and **Greg Durrett**. 2021. Modeling Fine-Grained Entity Types with Box Embeddings. *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)*.
49. Shih-ting Lin, Nathanael Chambers, and **Greg Durrett**. 2021. Conditional Generation of Temporally-ordered Event Sequences. *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)*.
48. Jiacheng Xu and **Greg Durrett**. 2021. Dissecting Generation Modes for Abstractive Summarization Models via Ablation and Attribution. *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)*.
47. Mahnaz Koupaee, **Greg Durrett**, Nathanael Chambers, and Niranjan Balasubramanian. 2021. Don’t Let Discourse Confine Your Model: Sequence Perturbations for Improved Event Language Models. *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)*.
46. Jifan Chen and **Greg Durrett**. 2021. Robust Question Answering Through Sub-part Alignment. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
45. Tanya Goyal and **Greg Durrett**. 2021. Annotating and Modeling Fine-grained Factuality in Summarization. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
44. Elisa Ferracane, **Greg Durrett**, Junyi Jessy Li and Katrin Erk. 2021. Did they answer? Subjective acts and intents in conversational discourse. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
43. Qiaochu Chen, Aaron Lamoreaux, Xinyu Wang, **Greg Durrett**, Osbert Bastani and Isil Dillig. 2021. Web Question Answering with Neurosymbolic Program Synthesis. *Proceedings of the Conference on Programming Language Design and Implementation (PLDI)*.
42. Shrey Desai, Jiacheng Xu, and **Greg Durrett**. 2020. Compressive Summarization with Plausibility and Salience Modeling. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
41. Jiacheng Xu, Shrey Desai, and **Greg Durrett**. 2020. Understanding Neural Abstractive Summarization Models via Uncertainty. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
40. Shrey Desai and **Greg Durrett**. 2020. Calibration of Pre-trained Transformers. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
39. Wei-Jen Ko, Te-Yuan Chen, Yiyang Huang, **Greg Durrett**, and Junyi Jessy Li. 2020. Inquisitive Question Generation for High Level Text Comprehension. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.

38. Tanya Goyal and **Greg Durrett**. 2020. Evaluating Factuality in Generation with Dependency-level Entailment. *Findings of the Conference on Empirical Methods for Natural Language Processing (Findings of EMNLP)*.
37. Yasumasa Onoe and **Greg Durrett**. 2020. Interpretable Entity Representations through Large-Scale Typing. *Findings of the Conference on Empirical Methods in Natural Language Processing (Findings of EMNLP)*.
36. Kaj Bostrom and **Greg Durrett**. 2020. Byte Pair Encoding is Suboptimal for Language Model Pretraining. *Findings of the Conference on Empirical Methods in Natural Language Processing (Findings of EMNLP)*.
35. Xi Ye, Qiaochu Chen, Xinyu Wang, Isil Dillig, and **Greg Durrett**. Sketch-Driven Regular Expression Generation from Natural Language and Examples. *Transactions of the Association for Computational Linguistics (TACL)*.
34. Tanya Goyal and **Greg Durrett**. 2020. Neural Syntactic Preordering for Controlled Paraphrase Generation. *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*.
33. Xi Ye, Qiaochu Chen, Isil Dillig, and **Greg Durrett**. 2020. Benchmarking Multimodal Regex Synthesis with Complex Structures. *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*.
32. Qiaochu Chen, Xinyu Wang, Xi Ye, **Greg Durrett**, and Isil Dillig. 2020. Multi-Modal Synthesis of Regular Expressions. *Proceedings of the Conference on Programming Language Design and Implementation (PLDI)*.
31. Jiayi Wei, Maruth Goyal, **Greg Durrett**, and Isil Dillig. 2020. LambdaNet: Probabilistic Type Inference using Graph Neural Networks. *Proceedings of the International Conference on Learning Representations (ICLR)*.
30. Yasumasa Onoe and **Greg Durrett**. 2020. Fine-Grained Entity Typing for Domain-Independent Entity Linking. *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*.
29. Jiacheng Xu and **Greg Durrett**. 2019. Neural Extractive Text Summarization with Syntactic Compression. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
28. Aditya Gupta and **Greg Durrett**. 2019. Effective Use of Transformer Networks for Entity Tracking. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
27. Su Wang, **Greg Durrett**, and Katrin Erk. 2019. Query-focused Scenario Construction. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
26. Tanya Goyal and **Greg Durrett**. 2019. Embedding time expressions for deep temporal ordering models. *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*.
25. Elisa Ferracane, **Greg Durrett**, Junyi Jessie Li, and Katrin Erk. 2019. Evaluating Discourse in Structured Text Representations. *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*.
24. Yasumasa Onoe and **Greg Durrett**. 2019. Learning to Denoise Distantly-Labeled Data for Entity Typing. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
23. Jifan Chen and **Greg Durrett**. 2019. Understanding dataset design choices for multi-hop reasoning. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
22. Wei-Jen Ko, **Greg Durrett**, and Junyi Jessie Li. 2019. Linguistically-Informed Specificity and Semantic Plausibility for Dialog Generation. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
21. Wei-Jen Ko, **Greg Durrett**, and Junyi Jessie Li. 2019. Domain Agnostic Real-Valued Specificity Prediction. *Proceedings of the AAAI Conference on Artificial Intelligence*.
20. Jiacheng Xu and **Greg Durrett**. 2018. Spherical Latent Spaces for Stable Variational Autoencoders. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
19. David Mueller and **Greg Durrett**. 2018. Effective Use of Context in Noisy Entity Linking. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
18. Su Wang, Eric Holgate, **Greg Durrett**, and Katrin Erk. 2018. Picking Apart Story Salads. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
17. Su Wang, **Greg Durrett**, and Katrin Erk. 2018. Modeling Semantic Plausibility by Injecting World Knowledge. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.

16. **Greg Durrett**, Jonathan Kummerfeld, Taylor Berg-Kirkpatrick, Rebecca Portnoff, Sadia Afroz, Damon McCoy, Kirill Levchenko, and Vern Paxson. 2017. Identifying Products in Online Cybercrime Marketplaces: A Dataset for Fine-grained Domain Adaptation. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*
15. Rebecca Portnoff, Sadia Afroz, **Greg Durrett**, Jonathan Kummerfeld, Taylor Berg-Kirkpatrick, Damon McCoy, Kirill Levchenko, and Vern Paxson. 2017. Tools for Automated Analysis of Cybercriminal Markets. *Proceedings of the International World Wide Web Conference (WWW)*.
14. **Greg Durrett**, Taylor Berg-Kirkpatrick, and Dan Klein. 2016. Learning-Based Single-Document Summarization with Compression and Anaphoricity Constraints. *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (ACL)*.
13. Matthew Francis-Landau, **Greg Durrett**, and Dan Klein. 2016. Capturing Semantic Similarity for Entity Linking with Convolutional Neural Networks. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
12. **Greg Durrett** and Dan Klein. 2015. Neural CRF Parsing. *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics (ACL)*.
11. James Ferguson, **Greg Durrett** and Dan Klein. 2015. Disfluency Detection with a Semi-Markov Model and Prosodic Features. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
10. **Greg Durrett** and Dan Klein. 2014. A Joint Model for Entity Analysis: Coreference, Typing, and Linking. *Transactions of the Association for Computational Linguistics (TACL)*.
9. David Hall, **Greg Durrett**, and Dan Klein. 2014. Less Grammar, More Features. *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics (ACL)*.
8. **Greg Durrett** and Dan Klein. 2013. Easy Victories and Uphill Battles in Coreference Resolution. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)* **Best Paper Finalist**.
7. **Greg Durrett**, David Hall, and Dan Klein. 2013. Decentralized Entity-Level Modeling for Coreference Resolution. *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (ACL)*.
6. Taylor Berg-Kirkpatrick, **Greg Durrett**, and Dan Klein. 2013. Unsupervised Transcription of Historical Documents. *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (ACL)*
5. **Greg Durrett** and John DeNero. 2013. Supervised Learning of Complete Morphological Paradigms. *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
4. **Greg Durrett**, Adam Pauls, and Dan Klein. 2012. Syntactic Transfer Using a Bilingual Lexicon. *Proceedings of the Conference on Empirical Methods for Natural Language Processing (EMNLP)*.
3. **Greg Durrett** and Dan Klein. 2011. An Empirical Investigation of Discounting in Cross-Domain Language Models. *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics (ACL)*
2. **Greg Durrett**, Frank Neumann, and Una-May O'Reilly. 2011. Computational Complexity Analysis of Simple Genetic Programming On Two Problems Modeling Isolated Program Semantics. *Foundations of Genetic Algorithms*.
1. **Greg Durrett**, Muriel Médard, and Una-May O'Reilly. 2010. A Genetic Algorithm to Minimize Chromatic Entropy. *Proceedings of the European Conference on Evolutionary Computation in Combinatorial Optimization (EvoCOP)*.

Accept rates for most NLP conferences can be found at: https://aclweb.org/aclwiki/Conference_acceptance_rates. *Findings* venues represent the next 15% of papers after the top 25% appearing in the main conference.

WORKSHOPS AND PREPRINTS

9. Aniruddh Sriram, Fangyuan Xu, Eunsol Choi, and **Greg Durrett**. 2024. Contrastive Learning to Improve Retrieval for Real-world Fact Checking. *Seventh Workshop on Fact Extraction and Verification (FEVER) at EMNLP (oral presentation)*
8. Zayne Sprague, Kaj Bostrom, Swarat Chaudhuri, and **Greg Durrett**. 2023. Deductive Additivity for Planning of Natural Language Proofs. *Proceedings of the First Workshop on Natural Language Reasoning and Structured Explanations*.
7. Tanya Goyal, Junyi Jessy Li, and **Greg Durrett**. 2022. News Summarization and Evaluation in the Era of GPT-3. *arXiv*.

6. Liyan Tang, Dhruv Rajan, Suyash Mohan, Abhijeet Pradhan, R. Nick Bryan, and **Greg Durrett**. 2021. Making Document-Level Information Extraction Right for the Right Reasons. *arXiv*.
5. Nila Selvaraj, Yasumasa Onoe, and **Greg Durrett**. 2021. Cross-Lingual Fine-Grained Entity Typing. *arXiv*.
4. Xinyu Zhao, Shih-ting Lin, and **Greg Durrett**. 2021. Effective Distant Supervision for Temporal Relation Extraction. *Proceedings of Adapt-NLP: The Second Workshop on Domain Adaptation for NLP (at EACL) 2021*.
3. Su Wang, **Greg Durrett**, and Katrin Erk. 2020. Narrative Interpolation for Generating and Understanding Stories. *arXiv*.
2. Jifan Chen, Shih-ting Lin, and **Greg Durrett**. 2019. Multi-hop Question Answering via Reasoning Chains. *arXiv*.
1. Aditya Gupta and **Greg Durrett**. 2019. Tracking Discrete and Continuous Entity State for Process Understanding. *NAACL 2019 Workshop on Structured Prediction for NLP (oral)*.

GRANTS, AWARDS, AND FELLOWSHIPS

- Exxon Sponsored Research Project, May 2025–May 2026
- Outstanding Paper Award at EMNLP 2024, November 2024
- NSF-Simons AI Institute for Cosmic Origins (lead PI: Stella Offner), Oct. 2024–Sept 2029
- Creating Connections, Applied Research Labs / UT Austin joint program, Aug. 2023–May 2024
- NLM R01: Closing the loop with an automatic referral population and summarization system (Yifan Peng (primary), Justin Rousseau, Ying Ding), July 2023–June 2028
- Amazon Science Hub Research Award, June 2023
- Sony Research Award (with Eunsol Choi), March 2023
- **Sloan Fellowship, Feb. 2023**
- Gift from Adobe Research, Sept. 2022
- **NSF CAREER: Flexible and Robust Reasoning in Natural Language, July 2022–June 2027**
- Open Philanthropy Project on AI Alignment: Seeking Edits and Explanations from Users to Generate Long, Knowledge-Rich Texts (with Eunsol Choi, Ray Mooney), June 2022–May 2025
- Gift from Salesforce Research, Jan. 2022
- Creating Connections, Applied Research Labs / UT Austin joint program, April 2022–Aug. 2023
- Good Systems, Office of the Vice President for Research, UT Austin: Designing Responsible AI Technologies to Protect Information Integrity (with Matt Lease (primary), Dhiraj Murthy, Jessy Li, Jo Lukito, Maria De-Arteaga, Min Kyung Lee), Sept. 2021–Aug. 2027
- Amazon Research Award (Alexa Fairness in AI), 2021
- Salesforce Research Deep Learning Grant, 2020
- Walmart Labs, Sponsored Research Project, 2020
- DARPA KAIROS (with Niranjana Balasubramanian (primary), Minh Hoai Nguyen, Frank Ferraro, Nate Chambers, Katrin Erk, Ray Mooney), Sept. 2019–Aug. 2023
- Gift from Arm, Inc., 2019
- NSF SaTC CORE: Towards Transparent and Effective Program Analysis for Detecting Mobile Malware (with Isil Dillig (primary), Yu Feng), 2019–2022
- Texas Health Catalyst Award, 2019
- **NSF Small: Handholding Supervision for Latent Variable Models in NLP (sole author), 2018–2021**
- NSF Medium: Computer-Aided Programming for Data Science (with Isil Dillig (primary), Ruben Martins), 2018–2021
- Bloomberg Data Science Research Grant, 2017–2018
- Facebook Fellowship in Natural Language Processing, 2014–2015
- Chair’s Special Award for service to the Berkeley EECS Department, May 2014
- Best Paper Finalist at EMNLP 2013, October 2013
- Outstanding Graduate Student Instructor Award, April 2012
- National Science Foundation Graduate Research Fellowship, April 2010
- National Merit Scholar, June 2006
- Robert C. Byrd Scholar, June 2006

TEACHING

Materials: Teaching materials (slides and course materials) from the following classes are publicly released with the following publication:

Greg Durrett, Jifan Chen, Shrey Desai, Tanya Goyal, Lucas Kabel, Yasumasa Onoe and Jiacheng Xu. 2021. Contemporary NLP Modeling in Six Comprehensive Programming Assignments *Proceedings of the Fifth Workshop on Teaching NLP*.

Select materials used at UC Berkeley, University of Washington, Carnegie Mellon University, Princeton, Columbia, Georgia Tech, University of Pennsylvania, Ohio State University, Duke University, ETH Zurich, George Mason University, UC Santa Barbara, and Penn State.

All materials, including 100+ video lectures available at <https://www.cs.utexas.edu/~gdurrett/>

Courses taught:

- **UT Austin CS388: Natural Language Processing (graduate)**
Instructor, UT Austin, January 2024 – May 2024, January 2023 – May 2023, January 2021 – May 2021; September 2019 – December 2019; August 2018 – December 2018.
- **UT Austin CS378: Natural Language Processing (undergraduate)**
Instructor, UT Austin, September 2024 – December 2024, September 2023 – December 2023, September 2022 – December 2022, September 2021 – December 2021, September 2020 – December 2020; January 2020 – May 2020; January 2019 – May 2019
- **UT Austin CS388: Natural Language Processing (Masters of Computer Science Online program)**
Instructor, UT Austin, September 2024 – December 2024, September 2023 – December 2023, September 2022 – December 2022, September 2021 – December 2021, January 2021 – May 2021
Note: this is a large class (800 students in fall 2024) and is taught in addition to my normal teaching load.
Press: <https://www.nytimes.com/2023/01/26/technology/ai-masters-degree-texas.html>
- **UT Austin CS395T: Structured Models for NLP (graduate)**
Instructor, UT Austin, August 2017 – December 2017
- **UC Berkeley CS288: Statistical Natural Language Processing (graduate)**
Teaching Assistant, UC Berkeley, August 2014 – December 2014
- **UC Berkeley CS188: Artificial Intelligence (undergraduate)**
Teaching Assistant, UC Berkeley, August 2011 – December 2011

High school outreach: Designed and recorded content for a module for use in high school classrooms.

Available online at <https://www.cs.utexas.edu/~gdurrett/courses/nlp-module/>

Press: <https://tacc.utexas.edu/news/latest-news/2025/02/06/ai-symposium/>

INVITED TALKS

- UMass Lowell, April 11, 2025
- New York University, March 3, 2025
- Yale, February 27, 2025
- University of Pennsylvania, February 25, 2025
- Oracle Distinguished Lecture Series, February 20, 2025
- University of Chicago, February 13, 2025
- Toyota Technological Institute at Chicago, February 11, 2025
- UMass Amherst NLP Seminar, November 20, 2024
- University of Chicago Communication & Intelligence Conference, October 18, 2024
- Cambridge University Language Technology Lab Seminar, October 17, 2024
- Snowflake ML, October 10, 2024
- Knowledge-Augmented NLP Workshop at ACL 2024, August 16, 2024
- *SEM workshop at NAACL 2024 (talk and panel discussion), June 21, 2024
- Panelist, “Generative AI”, Machine Learning Lab Research Symposium, April 18, 2024
- Georgetown University, November 28, 2023
- George Mason University, November 27, 2023
- University of Pennsylvania, November 13, 2023
- Princeton University, November 9, 2023
- Essentials of AI for Life and Society, UT Austin, October 11, 2023
- Texas A&M, Kingsville, October 4, 2023
- DialDoc Workshop at ACL 2023, July 13, 2023
- AMIA Journal Club, June 13, 2023
- Microsoft Research Cognitive Services Group, March 9, 2023
- Panelist, “ChatGPT and the Future of Healthcare”, IC² Institute, February 24, 2023
- NSF Expedition “Understanding the World Through Code” Webinar talk, January 31, 2023
- Amazon, January 25, 2023
- Stanford, October 20, 2022

- Workshop on Automatic Summarization of Creative Writing at COLING 2022, October 17, 2022
- Toyota Technological Institute at Chicago, September 29, 2022
- Salesforce, September 14, 2022
- Tencent, August 17, 2022
- University of Edinburgh, June 24, 2022
- Workshop on Commonsense Representation and Reasoning at ACL, May 27, 2022
- University of Washington, May 12, 2022
- Allen Institute for Artificial Intelligence, May 11, 2022
- Workshop in Explainable AI in Health at WWW 2022, April 25, 2022
- Baidu, April 19, 2022
- Cornell Tech, April 15, 2022
- NYU Text-as-Data, April 14, 2022
- Massachusetts Institute of Technology, April 1, 2022
- University of Southern California, January 14, 2022
- (Panelist) Workshop on Commonsense Reasoning and Knowledge Bases (CSKB) at AKBC 2021
- Automated Knowledge Base Construction (AKBC), October 6, 2021
- University of Pennsylvania, March 1, 2021
- UT Austin Institute for Foundations of Machine Learning Seminar Series, February 8, 2021
- UC Santa Barbara, January 14, 2021
- Microsoft Research, December 9, 2020
- Georgia Tech, October 23, 2020
- Salesforce Research, July 21, 2020
- Carnegie Mellon University, September 27, 2019
- WeCNLP Invited Talk on Low-Resource NLP, September 6, 2019
- Arm, March 16, 2019
- UT Random Structures Seminar, February 27, 2019
- Bloomberg, July 10, 2018
- Toyota Technological Institute at Chicago, June 25, 2018
- ICES Babuska Forum, UT Austin, March 2, 2018
- Berkeley AI Research Seminar, October 21, 2016
- The University of Texas at Austin, April 14, 2016
- University of California, Irvine, March 16, 2016
- New York University, March 9, 2016
- Carnegie Mellon University, March 7, 2016
- Princeton University, March 1, 2016
- University of California, San Diego, February 22, 2016
- Google New York (NLP group), February 8, 2016
- Microsoft Research, February 4, 2016
- ICSI Berkeley, November 25, 2014
- Facebook, November 20, 2014
- Google Mountain View (SAFT group), February 27, 2014

CONFERENCE AND WORKSHOP ORGANIZATION

- **Demo Chair for ACL 2026**
- **NAACL Board Member**, January 2024–January 2026
- **Co-organizer** (one of six) for the NAACL 2025 4th Workshop on Knowledge-Augmentation for Language Models and NLP Methods
- **Co-organizer** (one of seven) for the ACL 2024 workshop on Natural Language Reasoning and Structured Explanations
- **Co-organizer** (one of six) for the ACL 2023 workshop on Natural Language Reasoning and Structured Explanations
- **Co-organizer** (one of eight) for the ACL 2021 workshop on Natural Language Processing for Programming
- **Co-organizer** (one of seven) for the EMNLP 2019 workshop on Deep Learning for Low-Resource NLP (DeepLo)
- **Best paper award committee member** for NAACL 2025
- **Publication Chair** for EMNLP 2021
- **Faculty Advisor** (one of two) for the NAACL 2019 Student Research Workshop
- **Senior Area Chair:**
 - 2025:** NAACL 2025 (Language Modeling), ACL 2025 (Language Modeling), EMNLP 2025 (Safety and Alignment in LLMs)
 - 2024:** ACL 2024 (Sentence-level Semantics), NAACL 2024 (Summarization)
 - 2022:** NAACL 2022 (Summarization)
- **Action Editor** for ACL Rolling Review (starting June 2021)
- **Action Editor** for TMLR (since June 2023)

- **Area Chair:**

2026: ICLR 2026

2025: ICLR 2025, COLM 2025, NeurIPS 2025 (Datasets & Benchmarks)

2024: ICLR 2024, COLM 2024, NeurIPS 2024

2023: EMNLP 2023 (Interpretability), NeurIPS Datasets & Benchmarks 2023, NeurIPS 2023, ACL 2023 (Generation)

2022: NeurIPS Datasets & Benchmarks 2022, EMNLP 2022 (Information Extraction)

2021: AKBC 2021 (Information Extraction), ACL 2021 (Sentence-level Semantics), NAACL 2021 (Language Generation),

2020: CoNLL 2020 (Applications and Resources), EMNLP 2020 (Information Extraction)

2019: CoNLL 2019 (Tagging/Chunking/Syntax/Parsing)

- **Submission mentor** for the ACL 2023 Student Research Workshop (2 papers), NAACL 2022 SRW (2 papers), ACL 2022 SRW (1 paper) ACL 2021 SRW (1 paper), NAACL 2021 SRW (1 paper), ACL 2020 SRW (2 papers), ACL 2019 SRW (1 paper), NAACL 2019 SRW (1 paper)

- **Program Committee member (reviewer):**

2023: Big Picture Workshop at EMNLP 2023, ICML 2023

2022: NeurIPS 2022, ICML 2022

2021: New Frontiers in Summarization 2021, MRQA Workshop 2021, EMNLP 2021, NeurIPS 2021, ICLR 2021

2020: NeurIPS 2020, AKBC 2020, ICML 2020, ACL 2020

2019: EMNLP 2019, NeurIPS 2019, ACL 2019, ICML 2019, IEEE Sec. and Privacy 2019 (ad-hoc), NAACL 2019

2018: EMNLP 2018, ACL 2018, ICML 2018, GenDeep 2018, NAACL 2018, AAAI 2018 (ad-hoc)

2017: EMNLP 2017, New Frontiers in Summarization 2017, CoNLL 2017, ACL 2017, ACL 2017 Demo

2016: EMNLP 2016, COLING 2016, CoNLL 2016, ACL 2016, NAACL 2016

2015: EMNLP 2015, ACL 2015, NAACL 2015

2014: EMNLP 2014, ACL 2014, STAR-SEM 2014

2013: EMNLP 2013

- **Outstanding reviewer award** for NeurIPS 2022 (top 12%), EMNLP 2021, ICLR 2021, ICML 2020 (top 33% of reviewers), NeurIPS 2019 (top 50% of reviewers), EMNLP 2019 (top 1/6 of reviewers), EMNLP 2018 (top 7% / top 100 reviewers), ACL 2018, NAACL 2015 (top 25% of reviewers)
- **Expert Reviewer** for TMLR (June 2023)
- **Regular reviewer** for TMLR, 2022 – 2023
- **Regular reviewer** for TACL **since 2014** (>25 papers reviewed)
- Other journal reviewing: CL (1 paper), JAIR (2 papers (2023, 2021)), Computer Speech & Language (1 paper)

OTHER PROFESSIONAL SERVICE

- Reviewer for Schmidt Sciences AI2050 Proposal Review, May 2025
- CRA CCC Workshop “Future of Information Retrieval Research in the Age of Gen AI”; July 2024 (in-person)
- ISAT/DARPA The Great Oncoming kNowledge Ending (GONE) Workshop; March 2024 (in-person)
- Panelist for NSF Molecule Maker Lab Institute (MMLI) Site Visit (2 visits; in-person, University of Illinois, June 2023, July 2022)
- Panelist for NSF: 2025 (virtual panel), 2023 (virtual panel), 2018 (in-person panel)
- Ad hoc reviewer for U.S.-Israel Binational Science Foundation (1 proposal 2025)
- Ad hoc reviewer for Israel Science Foundation (1 proposal 2025, 1 proposal 2022, 1 proposal 2021)
- Ad hoc reviewer for European Research Council (1 proposal)
- Ad hoc reviewer for NSERC (1 proposal)
- Mentoring group leader: EMNLP 2021, NAACL 2021, ACL 2020, NAACL 2019
- Discussion group / Birds-of-a-Feather organizer: ACL 2021, NAACL 2021, ACL 2020, NAACL 2019)
- Conference session chair: NAACL 2022, ACL 2022, NAACL 2021, EMNLP 2020, ACL 2015

PHD THESES SUPERVISED

- Kaj Bostrom (PhD, 2024): Equipping Language Models for Systematic Reasoning
- Xi Ye (PhD, 2024): Steering Textual Reasoning with Explanations
→ Asst. Prof at University of Alberta after postdoc at Princeton
- Jocelyn (Qiaochu) Chen (PhD, 2024, co-supervised, primary supervisor: Isil Dillig): Neurosymbolic Program Synthesis for Information Extraction
→ Asst. Prof at University of Alberta
- Tanya Goyal (PhD, 2023): Fine-grained Evaluation for Text Summarization
→ Asst. Prof at Cornell University after postdoc at Princeton

- Jiayi Wei (PhD, 2023, co-supervised, primary supervisor: Isil Dillig): Combining Static Analysis with Deep Learning for Type Inference and Code Editing
→ AugmentCode
- Yasumasa Onoe (PhD, 2023, co-supervised with Eunsol Choi): Robust, Interpretable, and Usable Entity Representations
→ Google
- Jifan Chen (PhD, 2023, co-supervised with Eunsol Choi): Building Robust and Modular Question Answering Systems
→ Amazon AWS
- Wei-Jen Ko (PhD, 2022, co-supervised, primary supervisor: Jessie Li): Specific response generation and elaborative discourse structure
→ YouTube
- Jiacheng Xu (PhD, 2022): Neural Text Summarization with Fine-grained Control
→ Salesforce AI
- Su Wang (PhD, 2020, co-supervised, primary supervisor: Katrin Erk): Exploring Distributional Semantics in Lexical Representations and Narrative Modeling
→ Google

OTHER PHD COMMITTEES AT UT AUSTIN

- Yeming Wen (PhD 2025; supervisor: Swarat Chaudhuri)
- Ben Mariano (PhD 2024; supervisor: Isil Dillig)
- RJ Antonello (PhD 2024; supervisor: Alex Huth)
- Sheena Panthaplackel (PhD 2022; supervisor: Raymond Mooney)
- Jialin Wu (PhD 2022; supervisor: Raymond Mooney)
- Zhenpei Yang (PhD 2022; supervisor: Qixing Huang)
- Wei-Lin Hsiao (PhD 2021; supervisor: Kristen Grauman)
- Yanyao Shen (PhD 2020; supervisor: Sujay Sanghavi)
- Santiago Gonzalez (PhD 2020; supervisor: Risto Miikkulainen)
- Elisa Ferracane (PhD 2020; supervisors: Katrin Erk and Junyi Jessie Li)
- Pengxiang Cheng (PhD 2020; supervisor: Katrin Erk)
- Ye Zhang (PhD 2019; supervisor: Byron Wallace)
- Xinyu Wang (PhD 2019; supervisor: Isil Dillig)
- Elliott Meyerson (PhD 2018; supervisor: Risto Miikkulainen)
- Nazneen Rajani (PhD 2018; supervisor: Raymond Mooney)

EXTERNAL COMMITTEES

- Chaitanya Malaviya (PhD expected 2025, University of Pennsylvania; supervisor: Mark Yatskar, Dan Roth)
- Melanie Subbiah (PhD expected 2025, Columbia; supervisor: Kathy McKeown)
- Mahnaz Koupaee (PhD expected 2025, Stony Brook University; supervisor: Niranjan Balasubramanian)
- Yixin Liu (PhD expected 2025, Yale; supervisor: Arman Cohan)
- Darcey Riley (PhD expected 2025, University of Notre Dame; supervisor: David Chiang)
- Tomer Wolfson (PhD 2024, Tel Aviv University; supervisors: Jonathan Berant and Daniel Deutch): Democratizing Data Through Question Understanding
- Mathieu Ravaut (PhD 2024, NTU Singapore; supervisors: Sun Aixin and Joty Shafiq Rayhan): Neural Abstractive Summarization: Improvements at the Sequence-Level
- Sihao Chen (PhD 2024, University of Pennsylvania; supervisor: Dan Roth): Methods and Challenges in Inference Across Textual Sources
- Yilun Zhu (PhD 2023, Georgetown University; supervisor: Amir Zeldes): Advancing Generalization in Coreference Resolution
- Yumo Xu (PhD 2022, University of Edinburgh; supervisor: Mirella Lapata): Document Summarization with Neural Query Modeling
- Ali Emami (PhD 2021, McGill University; supervisor: Jackie Chi Kit Cheung): Towards the common-sense understanding of text: Results, resources and retrospection

UNDERGRADUATE AND MS THESES

- Prasann Singhal (BS, 2025; **CRA Outstanding Undergraduate Researcher Award Winner, Dean's Honored Graduate**): Understanding and Investigating Data Limitations in RLHF
→ PhD student at UC Berkeley
- Ramya Namuduri (BS, 2025): QUDSim: Quantifying Discourse Structure Similarities in LLM-Generated Text
- Anisha Gunjal (MS, 2024): Molecular Facts: Desiderata for Decontextualization in LLM Fact Verification
- Abhishek Divekar (MS, 2024): Diverse dataset synthesis using retrieval augmented language models

- Ritika Mangla (MS, 2024, co-supervised; primary supervisor: Jessy Li): Which questions should I answer? Saliency Prediction of Inquisitive Questions
- Aniruddh Sriram (BS, 2024): Improving Retrieval for Fact-Checking in the Wild
- Zayne Sprague (MS, 2023): Additive Deductive Embeddings for Planning of Natural Language Proofs
→ PhD student at UT Austin
- Jarad Forristal (BS, 2023): A Block Markov Chain Monte Carlo Sampling Method for Controllable LLM Inference
→ PhD student at UC San Diego
- Ryo Kamoi (MS, 2022): Shortcomings of Question Answering Based Factuality Frameworks for Error Localization
→ PhD student at Penn State University
- Ojas Ahuja (BS, 2022): Aspect-Oriented News Summarization
- Michael Royzen (BS, 2022): ASQA: A Dataset for Answer Sufficiency in QA
- Nivedhitha Selvaraj (BS, 2021; **Dean's Honored Graduate**): Cross-Lingual Fine-Grained Entity Typing
- Shih-Ting Lin (MS, 2020): Conditional Generation of Temporally-Ordered Event Sequences
- Shrey Desai (BS, 2020; **CRA Honorable Mention, Dean's Honored Graduate**): Compressive Summarization with Plausibility and Saliency Modeling
- Xinyu (Lucy) Zhao (BS, 2020): Learning Temporal Relations Across Domains
→ Jump Trading, then PhD student at MIT
- Akshay Kumar Gupta (BS, 2020): Topic-Focused Extractive Summarization
- Victoria Lestari (MS, 2018): Building Effective Representations for Domain Adaptation in Coreference Resolution

OTHER PROJECT SUPERVISION

- Grace (Ye Eun) Kim (BS, 2025) → PhD student at University of Pennsylvania
- Shankar Padmanabhan (BS, 2024; **CRA Honorable Mention, Dean's Honored Graduate**; primary supervisor: Eunsol Choi) → PhD student at Cornell University
- Adithya Bhaskar (summer intern, 2022) → PhD student at Princeton University
- Liyan Tang (BS, 2021) → PhD student at UT Austin
- Isabel Cachola (BS, 2019; **Dean's Honored Graduate**; primary supervisor: Jessy Li) → PhD student at Johns Hopkins University
- Dhruv Rajan (BS, 2019)
- Aditya Gupta (MS, 2019)
- David Mueller (BS, 2018) → PhD student at Johns Hopkins University

DEPARTMENTAL SERVICE

- Graduate Advisor for the Masters of Science in Data Science Online, June 2022-present
- Lead organizer for “Intro to Grad School” series with UT CS diversity student orgs, 2021
- Undergraduate Studies Committee (departmental), UT Austin Computer Science, 2020-2022
- Linguistics faculty hiring committee (computational position), UT Austin, 2020-2021
- Admissions Committee member, UT Austin, 2016-2017, 2017-2018, 2018-2019, 2019-2020
- Graduate course requirements committee member (departmental), 2017-2018
- Coordinator for the UT Austin AI retreat, Fall 2017
- Student coordinator for the 2014 UC Berkeley Computer Science Graduate Visit Weekend

INDUSTRY POSITIONS

- **Bespoke Labs**
Advisor, June 2024 –
- **Semantic Machines**, Berkeley, CA
Senior Research Scientist, June 2016 – June 2017
- **Google**, Mountain View, CA
Software Engineering Intern–Research, May 2012 – August 2012
- **Charles River Analytics**, Cambridge, MA
Software Engineering Intern, June 2008 – August 2008
- **ATC–NY**, Ithaca, NY
Intern, June 2007 – August 2007